

# Request Reconstruction in MirageOS Unikernels


Al Amjad Tawfiq Isstaif (aati2)

University of Cambridge

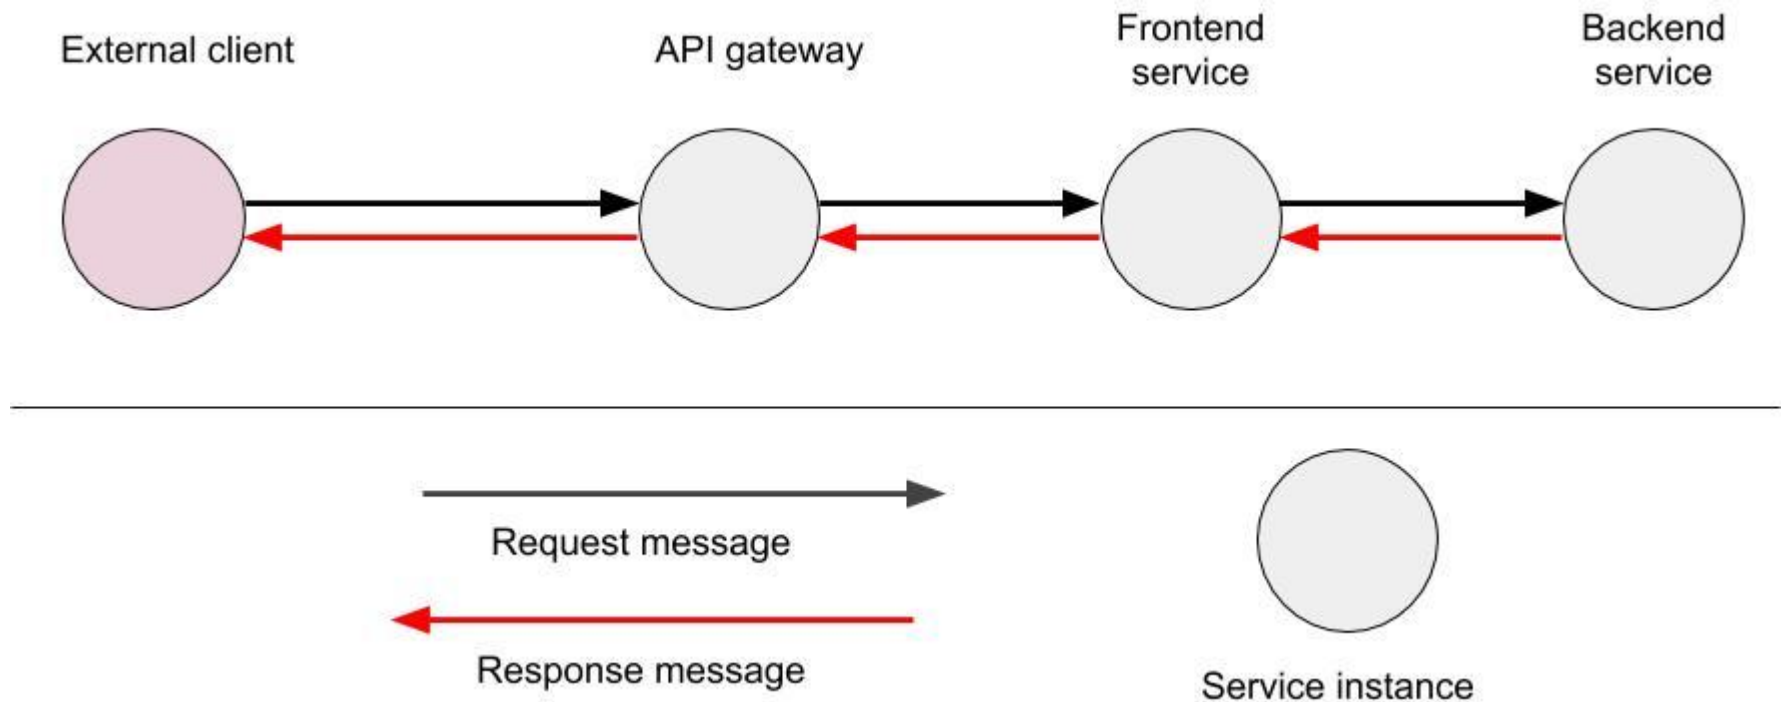
21 Jun 2019

# Motivation

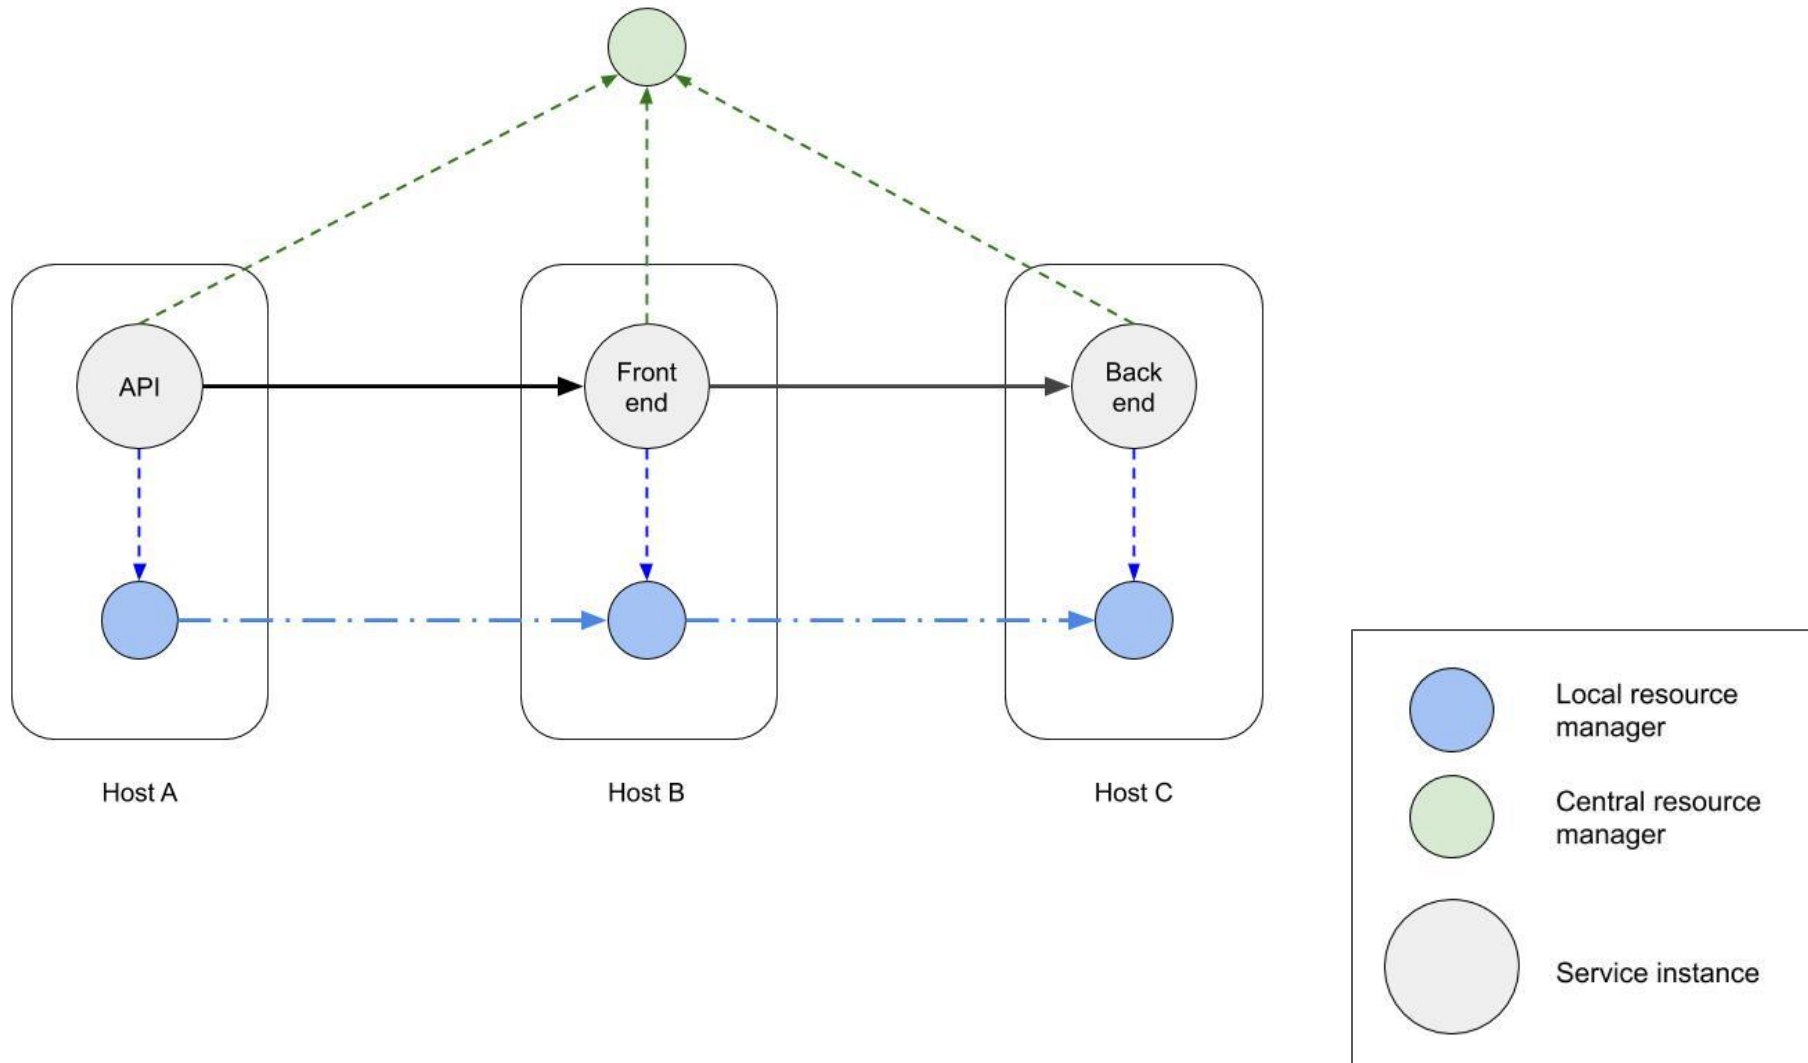
Exploring the potential of combining unikernels with distributed tracing to address the autoscaling problem in microservice applications



# Running example



# Trace collection and resource management



# Key ideas

Key ideas in the proposed tracing model:

- Measuring resource waiting times (CPU and Network)
- “Flexible” tracing analysis (application-level & service-level trace analysis)


Generic execution wrappers to instrument Mirage HTTP library (Cohttp)

- Protocol agnostic: can be used to automatically instrument any protocol libraries based on inter-service communication
- Mirage libraries based on Lwt lightweight threads or *promises (futures)*

API gateway

Frontend  
service

Backend  
service



Total: 50ms  
Local-wait: 1ms

Net-wait: 4ms

Total: 30ms  
Local-wait: 1ms

Net-wait: 5ms

Total: 20ms  
Local-wait: 10ms

Total: ...ms  
Local-wait: ...ms

Local context




Net-wait: ...ms


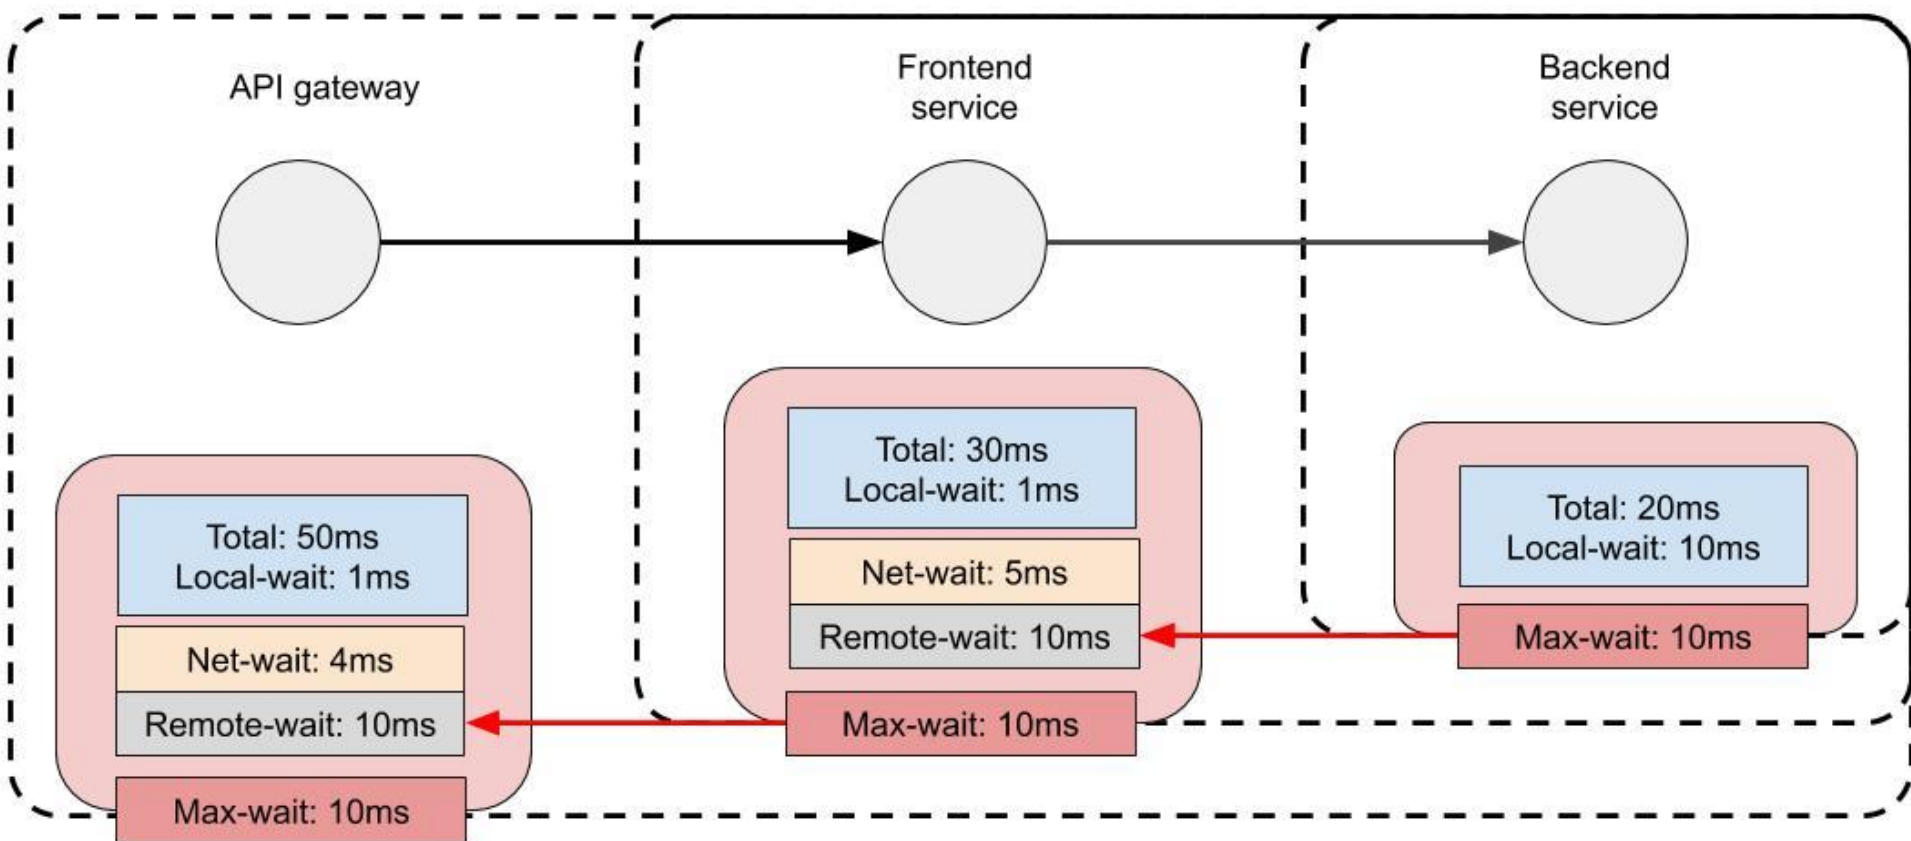
Remote context

API gateway


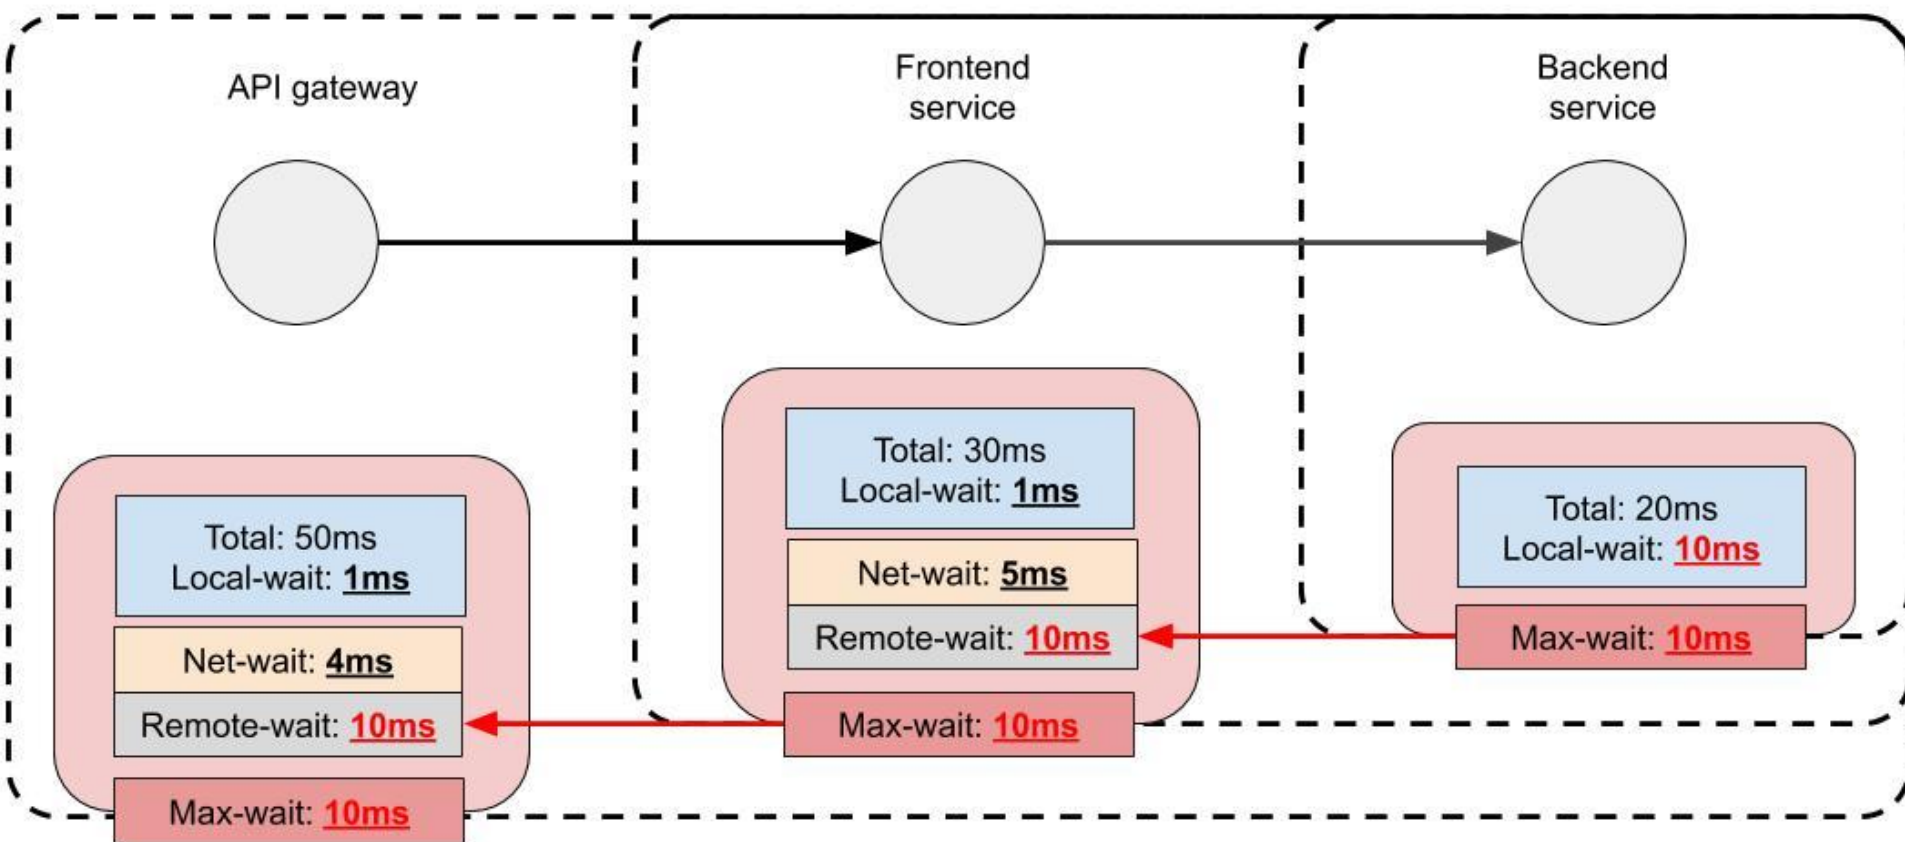
Frontend  
service

Backend  
service

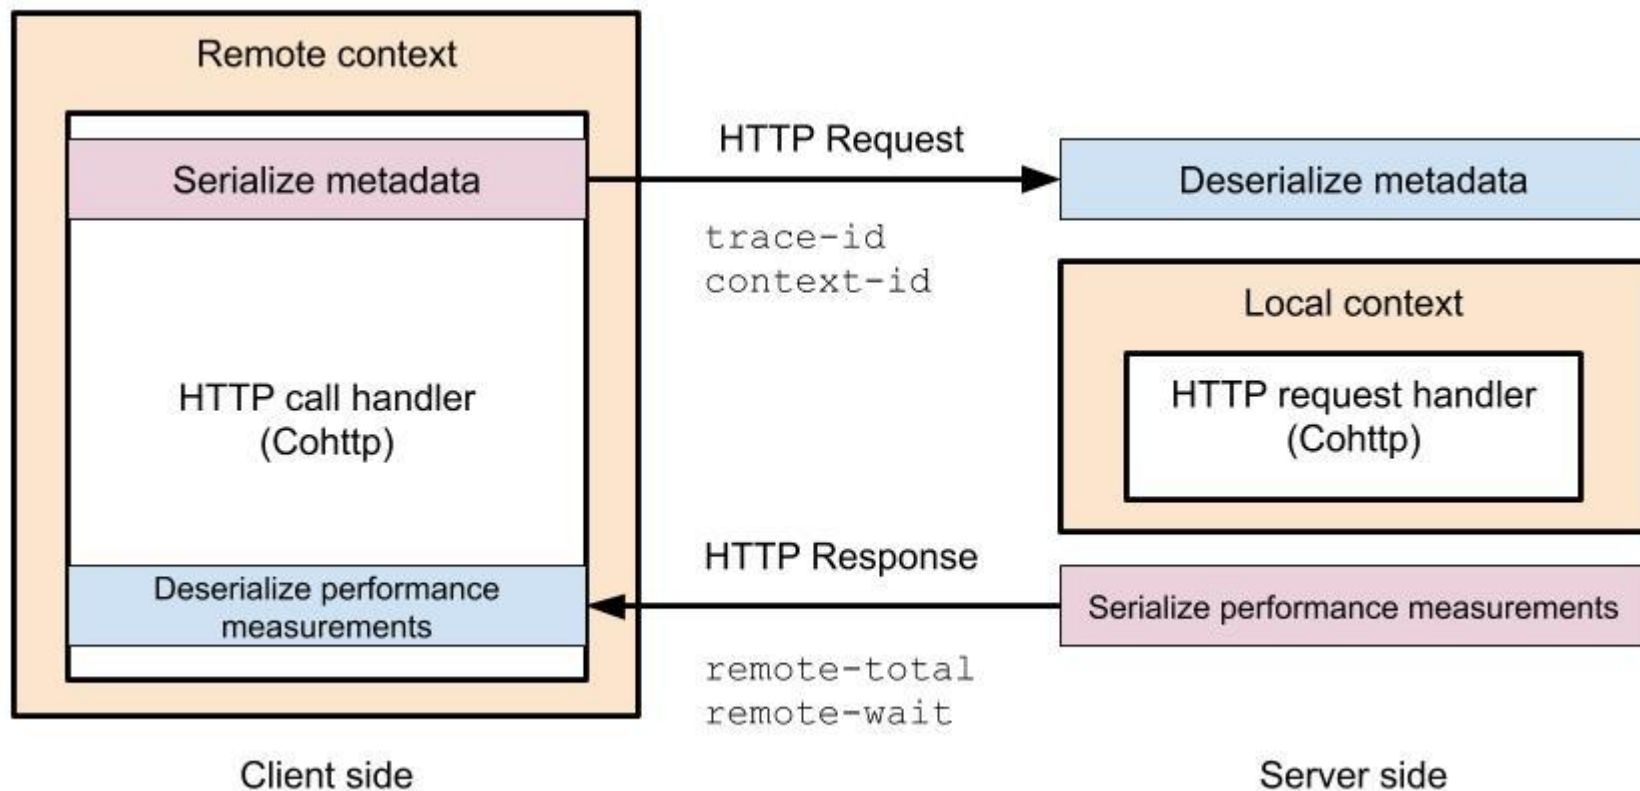









# Execution wrappers and Cohttp library instrumentation



# Overhead evaluation

For 100 explicit context switch points (also known as *yield points*)

- 8% overhead with no sampling
- Less than 2% under a sampling rate of 1/1024



# Summary and conclusions

A tracing model based on hierarchical aggregation of resource waiting times suitable for further research on microservices autoscaling policies

- The tracing model can be extended for other lightweight virtualization options (e.g. containers)


The single-threaded architecture of Mirage unikernels introduce various issues that need to be addressed in future work

- A CPU bottleneck can lead to an indistinguishable artificial network bottleneck
- Instrumenting the TCP stack can be useless, and packet-timestamping is required


Backup slides

# CPU bottleneck

Backend service




Frontend service




# Net bottleneck


Backend service



Frontend service



# Implementation of local context wrapper



```
let t = Lwt.with_context (fun () -> f)
```



```
let t = Lwt.with_context (fun () -> a >>= b >>= c)
```



Avoiding an artificial network bottleneck (above) by increasing the number of context switches (below)

